SLIDE 01 // COGNIFLOW AI

CogniFlow AI

Autonomous Multi-Agent Reliability & Enterprise QA Platform

- Category: Enterprise Software & GenAI Infrastructure
- Round: Series Seed | Raising $2.5M at $18.0M Post-Money Valuation
- Founders: Aarav Sharma (Ex-Postman) & Neha Verma (Stanford MS CS)

SLIDE 02 // THE HAIR-ON-FIRE PROBLEM

The Hair-On-Fire Problem

Enterprises Are Bleeding Money on Failed Production Agents

- Critical Pain Point: Enterprise LLM agents fail in production >45% of the time.
- Database Corruption: Non-deterministic agent hallucinated SQL writes corrupt production state.
- Wasted Capital: Fortune 500 enterprises burn over $300K/quarter on redundant token inference loops.

SLIDE 03 // THE 10X MECHANISM

The 10x Mechanism

Active Kernel-Level Interception with Zero-Downtime State Rollback

- Our Solution: A lightweight runtime observability SDK that intercepts agent tool calls in under 12ms.
- Tracks distributed execution trees and automatically rolls back corrupted database state.
- Zero Downtime: Eliminates 6 hours of average engineer manual rollback downtime per incident.

SLIDE 04 // WHY NOW?

Why Now?

The Macro Inflection Point in Multi-Agent Reasoning

- 1. Reasoning Models (o1/Sonnet/Gemini) just made autonomous multi-step agents viable in 2026.
- 2. Enterprise GenAI spending is shifting from prompt toys to compliance, reliability, and infrastructure.
- 3. eBPF kernel tracing maturity in Kubernetes enables near-zero latency traffic inspection.

SLIDE 05 // MARKET SIZING (BOTTOM-UP MATH)

Market Sizing (Bottom-Up Math)

From a $2.1B Beachhead to a $95.0B Global TAM

- Claimed TAM: $95.0B Global Enterprise Observability & AI Governance Software.
- Serviceable Obtainable Market (SOM): $2.16B (120,000 mid-to-large dev teams @ $18,000 ACV).
- Bottom-up math grounded in developer team volume and enterprise software pricing power.

SLIDE 06 // TECHNICAL MOAT & DEFENSIBILITY

Technical Moat & Defensibility

Why Foundation Models Cannot Replicate Our Infrastructure

- Kernel-Level eBPF Probes: Runs inside customer VPCs with zero latency overhead.
- Proprietary Local SLM Arbiter: Fine-tuned 3B local model replaces external API calls, cutting COGS by 76%.
- Runtime Lock-In: Deeply embedded in CI/CD pipeline test gates with high enterprise switching costs.

SLIDE 07 // COMPOUNDING TRACTION & ECONOMICS

Compounding Traction & Economics

Demonstrated Capital-Efficient Software Scale

- Financial Traction: $1,020,000 ARR ($85,000 MRR) growing at 5.7x YoY.
- Sales Efficiency: CAC is $4,800 with an elite 4.4 Months CAC Payback Period.
- Unit Economics: Inferred LTV:CAC of 25.28x with 91.0% Gross Margin and 0.9% monthly churn.
- Capital Discipline: Burn Multiple is 0.95x with 11.5 Months of Cash Runway ($750K in bank).

SLIDE 08 // GO-TO-MARKET FLYWHEEL

Go-To-Market Flywheel

Product-Led Developer Land & Enterprise Expand

- 1. Land: Free open-source Python/TypeScript agent SDK drives bottom-up adoption by engineering teams.
- 2. Expand: Usage-based pricing tiers charging per 100K monitored agent executions.
- 3. Lock: Enterprise air-gapped VPC deployments with SOC2 Type II, RBAC, and PII redaction.

SLIDE 09 // COMPETITOR LANDSCAPE (2X2)

Competitor Landscape (2x2)

Positioned at Active State Protection with Low Inference COGS

- Y-Axis: Active Runtime Rollback vs. Passive Log Ingestion.
- X-Axis: Lightweight Local SLM vs. Expensive Third-Party API Bills.
- Legacy APMs (Datadog/Dynatrace): High cost, passive logging only.
- OpenAI API Wrappers: High latency, 90% gross margin illusion, zero enterprise data privacy.

SLIDE 10 // FOUNDING TEAM & EARNED SECRETS

Founding Team & Earned Secrets

Systems Architects with Proven Execution Pedigree

- Aarav Sharma (CEO & Co-Founder): Ex-Staff Engineer at Postman. 2 US patents in distributed tracing.
- Neha Verma (CTO & Co-Founder): Ex-Machine Learning Lead at BrowserStack. M.S. CS from Stanford University.
- Advisors: Former VP of Infrastructure at Datadog and Lead Architect at HashiCorp.

SLIDE 11 // THE ASK & 18-MONTH MILESTONES

The Ask & 18-Month Milestones

Raising $2.5M Series Seed to Compound to $3.5M ARR

- Round Structure: Raising $2,500,000 Series Seed for 24 months of runway.
- Use of Funds: 70% Core Engineering (eBPF & distributed systems), 20% DevRel GTM, 10% GPU Cluster.
- Series A Milestones: Scale from $1.02M ARR to $3.5M ARR with 150+ Enterprise Accounts.