SLIDE 01 // CATEGORY HOOK

CogniFlow AI

Autonomous Multi-Agent Reliability & Enterprise QA Platform

Target Category Enterprise Software / GenAI Infrastructure	**Current Round** **Raising $2.5M Series Seed**	**Post-Money Valuation** $18.0M Post-Money Cap	**Headquarters** Bengaluru & San Francisco

The One-Liner: CogniFlow intercepts non-deterministic agent tool calls at the Linux kernel level in under 12ms, eliminating database corruption and automated hallucination cascades for Fortune 500 engineering teams.

SLIDE 02 // THE HAIR-ON-FIRE PROBLEM

Enterprises Are Bleeding Money on Broken Agent Loops

Critical Pain Point: Multi-agent workflows are failing in production due to 3 catastrophic vulnerabilities:

45% EXECUTION FAILURE RATE	DATABASE CORRUPTION	$300K+ WASTED TOKENS / QTR
Multi-agent loops get trapped in circular reasoning calls, causing enterprise web workers to hang indefinitely.	Hallucinated SQL and API writes corrupt production customer data, requiring an average of 6 hours of dev rollback downtime.	Uncontrolled recursive LLM tool calls burn enterprise cloud budgets on redundant API pass-through calls.

SLIDE 03 // THE 10x MECHANISM

Active Runtime Interception & Zero-Downtime State Rollback

WITHOUT COGNIFLOW (THE STATUS QUO)

- Agent initiates uncontrolled SQL / API mutation
- Database corrupts user account balances or records
- Incident triggers on PagerDuty 4 hours later
- Engineering team manually runs database restores
- **Average Loss: $65,000 per incident**

WITH COGNIFLOW (THE 10X MECHANISM)

- **Kernel-Level eBPF Probe intercepts tool call in <12ms**
- Local SLM Arbiter checks execution tree determinism
- Hallucinated state change is isolated & rolled back instantly
- Zero customer impact; full telemetry trace logged in CI/CD
- **Zero downtime & 100% compliance SLA**

SLIDE 04 // MACRO INFLECTION POINT

Why This Couldn't Exist 3 Years Ago

Three structural tailwinds make 2026-2027 the exact window to capture this market:

1. REASONING MODELS ARRIVED	2. BUDGET INFLECTION POINT	3. eBPF KERNEL MATURITY
Multi-step reasoning models (o1, Sonnet, Gemini) just made autonomous agentic loops viable in enterprise production.	Enterprise GenAI budgets have moved from 'exploratory prompt experiments' to 'production infrastructure & compliance SLAs'.	Linux kernel tracing probes now run inside enterprise Kubernetes clusters with near-zero latency overhead (<12ms).

SLIDE 05 // BOTTOM-UP MARKET SIZING

A $2.1B Beachhead Expanding to a $95B Global TAM

SOM (BEACHHEAD MARKET)	SAM (SERVICEABLE ADDRESSABLE)	TAM (TOTAL GLOBAL MARKET)
$2.16 BILLION	**$8.64 BILLION**	**$95.0 BILLION**
120,000 Mid-to-Large Engineering Teams × $18,000 Initial ACV. Target: US/EU SaaS and Indian Global Capability Centers (GCCs).	Fortune 2000 Enterprises deploying multi-agent workflows across production databases with dedicated VPC air-gaps.	Total Enterprise Observability, Automated QA, and AI Governance Software by 2030 (Gartner / IDC forecast).

SLIDE 06 // TECHNICAL DEFENSIBILITY

Why Foundation Models & Wrappers Can't Copy Us

1. KERNEL-LEVEL eBPF PROBES

Unlike Python wrapper tools (LangSmith, basic logging), our interceptor runs as an eBPF probe at the Linux OS layer inside customer VPCs. Zero network latency overhead.

2. LOCAL SLM ARBITER

Replaced expensive $30K/mo OpenAI evaluation bills with fine-tuned 3B parameter local models running on customer hardware, **slashing token COGS by 76%**.

3. RUNTIME CI/CD LOCK-IN

Embedded directly into developer testing pipelines and deployment gates. Ripping out CogniFlow creates immediate enterprise compliance vulnerabilities.

SLIDE 07 // TRACTION & COMPOUNDING ECONOMICS

Demonstrated Capital-Efficient Software Scale

ANNUAL RECURRING REVENUE	LTV : CAC RATIO	CAC PAYBACK PERIOD	BURN MULTIPLE
$1,020,000 ARR	**25.3x**	**4.4 MONTHS**	**0.95x**
$85,000 MRR | 5.7x YoY Growth	Benchmark: > 3.0x (Elite)	Fast cash flow recycling	David Sacks Metric: < 1.0x

Customer Health: 91.0% Software Gross Margins, 0.9% Monthly Logo Churn, and 11.5 Months of existing cash runway ($750K in bank).

SLIDE 08 // GO-TO-MARKET FLYWHEEL

Developer Land & Enterprise Expand Engine

PHASE 1: LAND (PLG)	PHASE 2: EXPAND (USAGE)	PHASE 3: LOCK (ENTERPRISE)
Free open-source Python and TypeScript agent SDK drives grassroots adoption among developers on GitHub.	Automated transition to usage-based pricing ($0.05 per 100K monitored agent executions) as applications enter production.	Annual Enterprise Contracts ($25K - $100K ACV) for air-gapped VPC deployments, SOC2 Type II, RBAC, and PII redaction.

SLIDE 09 // COMPETITIVE LANDSCAPE

Positioned at Active State Protection with Low Inference COGS

COMPETITIVE QUADRANT	HIGH INFERENCE COGS (API CALLS)	LOW INFERENCE COGS (LOCAL SLM)
ACTIVE RUNTIME ROLLBACK	**API Wrappers** High latency, 90% margin illusion, breaks when OpenAI changes API.	**COGNIFLOW AI (LEADER)** Kernel eBPF interception (<12ms), zero-token-overhead SLM, active rollback.
PASSIVE LOGGING ONLY	**Legacy APMs (Datadog/Dynatrace)** Very high cost, passive logs only; cannot rollback state.	**Open Source Logging (Prometheus)** Raw metrics, zero agent understanding, manual debugging.

SLIDE 10 // UNFAIR ADVANTAGE & TEAM

Systems Architects with Deep Execution Pedigree

AARAV SHARMA — *CEO & Co-Founder*

- **Ex-Staff Engineer @ Postman**: Built distributed tracing pipelines handling 25M developers.
- **Patents**: Holds 2 granted US patents in runtime execution interception.
- **Education**: B.Tech Computer Science (IIT Bombay).

NEHA VERMA — *CTO & Co-Founder*

- **Ex-ML Lead @ BrowserStack**: Architected distributed testing clusters running 100K jobs/hour.
- **Specialization**: Published researcher in small language model fine-tuning.
- **Education**: M.S. Computer Science (Stanford University).

SLIDE 11 // THE ASK & MILESTONES

Raising $2,500,000 to Scale to $3.5M ARR

CAPITAL ALLOCATION (USE OF FUNDS)

- **70% Core Distributed Systems Engineering**: Hiring senior eBPF kernel engineers in Bengaluru.
- **20% Developer Go-To-Market**: Technical DevRel, developer hackathons, and US enterprise sales.
- **10% Dedicated GPU Cluster**: Local SLM model fine-tuning and benchmark clusters.

18-MONTH SERIES A MILESTONES

- **ARR Target**: Scale from **$1.02M ARR to $3.5M ARR**.
- **Enterprise Logos**: Expand from 71 to **150+ Enterprise Accounts**.
- **Net Retention**: Maintain >130% Net Dollar Retention (NDR).
- **Series A Ready**: Positioned for a $12M - $15M round at top-decile valuation.